

Solving Problems by Searching

Two Types of Searching

- **Uninformed search** algorithms—algorithms that are given no information about the problem other than its definition
 - Examples: BFS, DFS, Uniform-Cost Search, Depth-Limited Search, Iterative deepening depth-first search, Bidirectional Search, etc.
- **Informed search** algorithms, on the other hand, can do quite well given some guidance on where to look for solutions
 - Examples: Greedy BFS, A* Search, iterative-deepening A* (IDA*), RBFS (recursive best-first search), SMA* (simplified memory-bounded A*), etc.

Problem Solving Agent

function SIMPLE-PROBLEM-SOLVING-AGENT(*percept*) **returns** an action
persistent:

seq, an action sequence, initially empty

state, some description of the current world state

goal, a goal, initially null

problem, a problem formulation

state \leftarrow UPDATE-STATE(*state*, *percept*)

if *seq* is empty **then**

goal \leftarrow FORMULATE-GOAL(*state*)

problem \leftarrow FORMULATE-PROBLEM(*state*, *goal*)

seq \leftarrow SEARCH(*problem*)

if *seq* = *failure* **then return** a null action

action \leftarrow FIRST(*seq*)

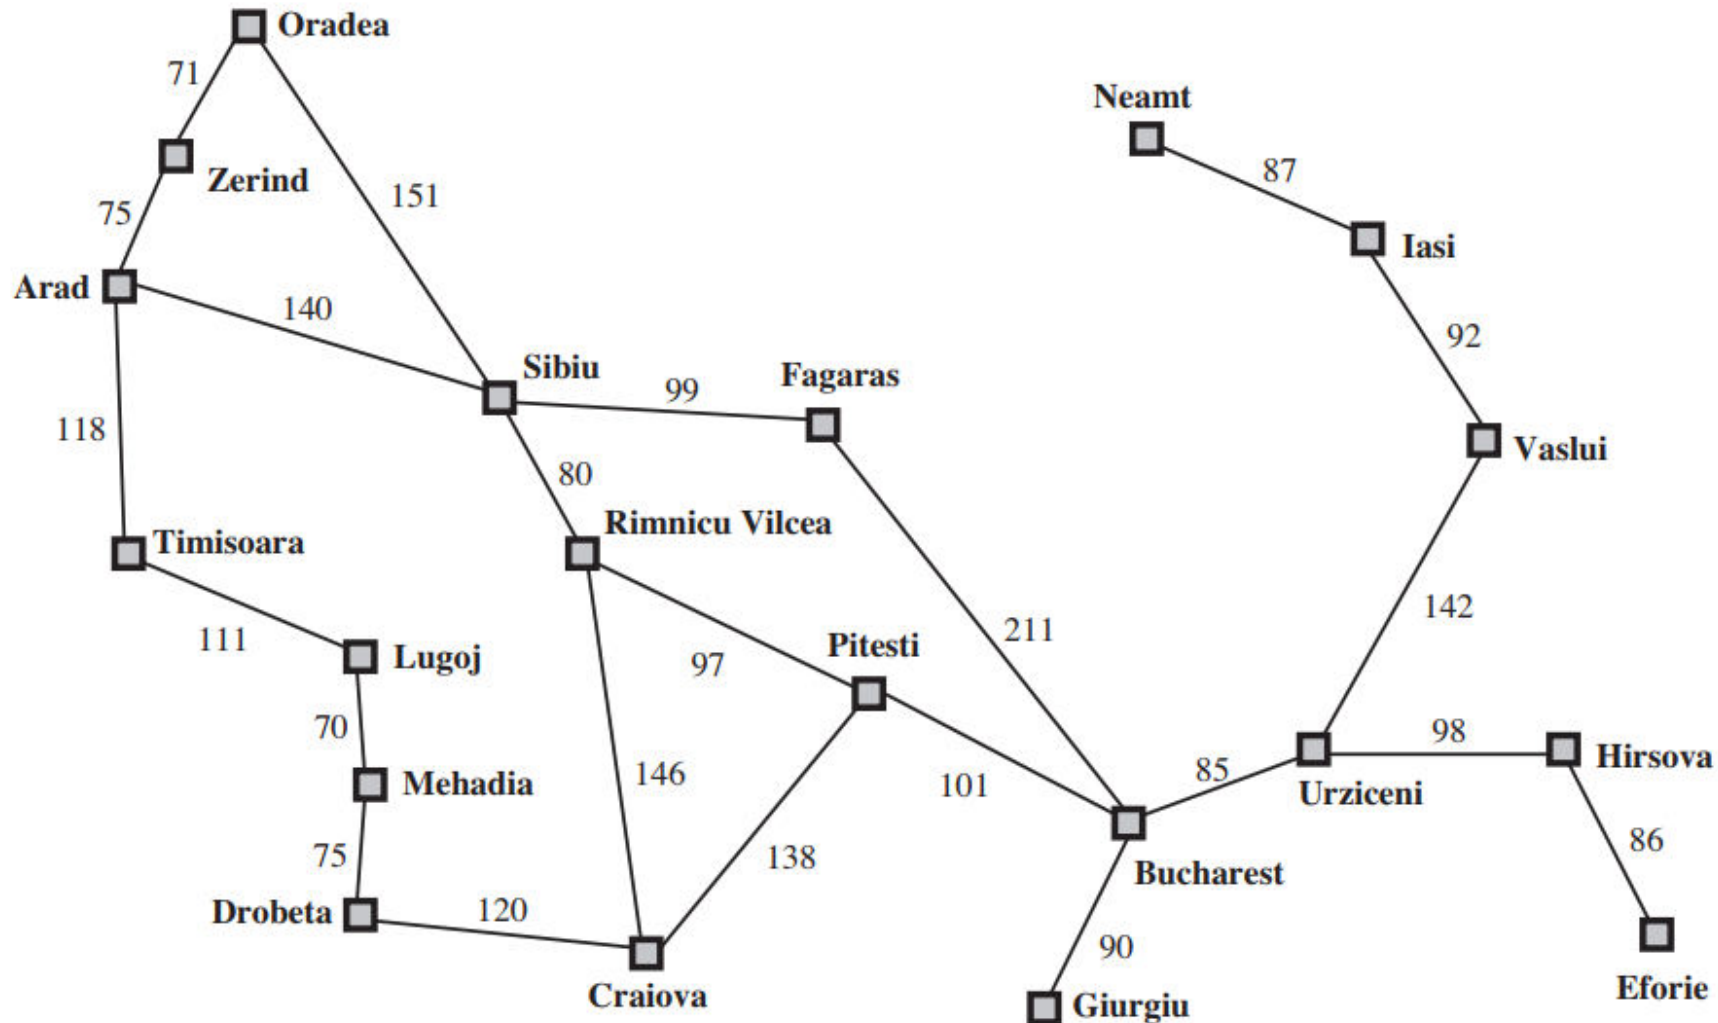
seq \leftarrow REST(*seq*)

return *action*

Problem Solving Agent

- **Goal formulation**, based on the current situation and the agent's performance measure, is the first step in problem solving.
- **Problem formulation** is the process of deciding what actions and states to consider, given a goal.
- The process of looking for a sequence of actions that reaches the goal is called **search**.
- A **search algorithm** takes a **problem** as input and returns a solution in the form of an **action sequence**.

A simplified road map of part of Romania



Problem Definition

- A problem can be defined formally by five components:
- **Initial State:** The **initial state** that the agent starts in.
 - For example, the initial state for our agent in Romania might be described as *In(Arad)*.
- **Action:** A description of the possible **actions** available to the agent.
 - Given a particular state s , $ACTIONS(s)$ returns the set of actions that can be executed in s .
 - For example, from the state *In(Arad)*, the applicable actions are $\{Go(Sibiu), Go(Timisoara), Go(Zerind)\}$.

Problem Definition

- **Transition Model:** A description of what each action does is the **transition model**
 - specified by a function $RESULT(s, a)$ that returns the state that results from doing action a in state s .
 - the term successor to refer to any state reachable from a given state by a single action.
 - For example, $RESULT(In(Arad), Go(Zerind)) = In(Zerind)$.
- Together, the initial state, actions, and transition model implicitly define the **state space** of the problem—the set of all states reachable from the initial state by any sequence of actions
- The state space forms a directed network or **graph** in which the nodes are states and the links between nodes are actions.
- A **path** in the state space is a sequence of states connected by a sequence of actions.

Problem Definition

- **Goal Test:** The **goal test**, which determines whether a given state is a goal state.
 - Sometimes there is an explicit set of possible goal states, and the test simply checks whether the given state is one of them.
 - The agent's goal in Romania is the singleton set $\{In(Bucharest)\}$.
- **Path Cost:** A **path cost** function that assigns a numeric cost to each path.
 - The problem-solving agent chooses a cost function that reflects its own performance measure.
 - For the agent trying to get to Bucharest, the cost of a path might be its length in kilometers.

Vacuum World

- This can be formulated as a problem as follows:
 - **States:** The state is determined by both the agent location and the dirt locations.
 - The agent is in one of two locations, each of which might or might not contain dirt. Thus, there are $2 \times 2^2 = 8$ possible world states. A larger environment with n locations has $n \times 2^n$ states.
 - **Initial state:** Any state can be designated as the initial state.
 - **Actions:** In this simple environment, each state has just three actions: Left, Right, and Suck. Larger environments might also include Up and Down.

Vacuum World

- This can be formulated as a problem as follows:
 - **Transition model:** The actions have their expected effects, except that moving Left in the leftmost square, moving Right in the rightmost square, and Sucking in a clean square have no effect.
 - **Goal test:** This checks whether all the squares are clean.
 - **Path cost:** Each step costs 1, so the path cost is the number of steps in the path.

Vacuum World




Figure: The state space for the vacuum world. Links denote actions: L = Left, R = Right, S = Suck.

8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

8-puzzle

- The standard problem formulation is as follows:
 - **States:** A state description specifies the location of each of the eight tiles and the blank in one of the nine squares.
 - **Initial state:** Any state can be designated as the initial state.
 - Note that any given goal can be reached from exactly half of the possible initial states.
 - **Actions:** The simplest formulation defines the actions as movements of the blank space Left, Right, Up, or Down. Different subsets of these are possible depending on where the blank is.

8-puzzle

- The standard problem formulation is as follows:
 - **Transition model:** Given a state and action, this returns the resulting state.
 - For example, if we apply Left to the start state (See in Figure), the resulting state has the 5 and the blank switched.
 - **Goal test:** This checks whether the state matches the goal configuration.
 - **Path cost:** Each step costs 1, so the path cost is the number of steps in the path.

Solution

- A **solution** to a problem is an action sequence that leads from the initial state to a goal state.
- Solution quality is measured by the path cost function.
- An **optimal solution** has the lowest path cost among all solutions.

Solve a Problem

- The possible action sequences starting at the initial state form a **search tree** with the initial state at the root.
- The **branches** are actions and the **nodes** correspond to states in the state space of the problem.
- Then **expand** the current state; that is, apply each legal action to the current state, thereby generate a new set of states.
- The set of all **leaf nodes** available for expansion at any given point is called the **frontier**.
- The process of expanding nodes on the frontier continues until either a solution is found or there are no more states to expand.
- There may be **repeated state** in the search tree.

Infrastructure for search algorithms

- Search algorithms require a data structure to keep track of the search tree.
- For each node n of the tree, we have a structure that contains four components:
 - n .STATE: the state in the state space to which the node corresponds;
 - n .PARENT: the node in the search tree that generated this node;
 - n .ACTION: the action that was applied to the parent to generate the node;
 - n .PATH-COST: the cost, traditionally denoted by $g(n)$, of the path from the initial state to the node, as indicated by the parent pointers.
- The appropriate data structure for this is a **queue**: FIFO queue, LIFO queue and Priority queue

Measuring problem-solving performance

- We can evaluate an algorithm's performance in four ways:
 - **Completeness:** Is the algorithm guaranteed to find a solution when there is one?
 - **Optimality:** Does the strategy find the optimal solution?
 - **Time complexity:** How long does it take to find a solution?
 - **Space complexity:** How much memory is needed to perform the search?

Measuring problem-solving performance

- In AI, the graph is often represented implicitly by the initial state, actions, and transition model and is frequently infinite.
- For these reasons, complexity is expressed in terms of three quantities:
 - b , the branching factor or maximum number of successors of any node
 - d , the depth of the shallowest goal node (i.e., the number of steps along the path from the root);
 - m , the maximum length of any path in the state space.
- Time is often measured in terms of the number of nodes generated during the search.
- Space in terms of the maximum number of nodes stored in memory.

Measuring problem-solving performance

- To assess the effectiveness of a search algorithm, we can consider one of the following:
- **Search cost:** which typically depends on the time complexity but can also include a term for memory usage.
- **Total cost:** which combines the search cost and the path cost of the solution found.

Uninformed Search

- Uninformed search (also called blind search)
- The strategies have no additional information about states beyond that provided in the problem definition.
- All they can do is generate successors and distinguish a goal state from a non-goal state.
- All search strategies are distinguished by the order in which nodes are expanded.
- Examples: BFS, DFS, Uniform-Cost Search, Depth-Limited Search, Iterative deepening depth-first search, Bidirectional Search, etc (**self study**)

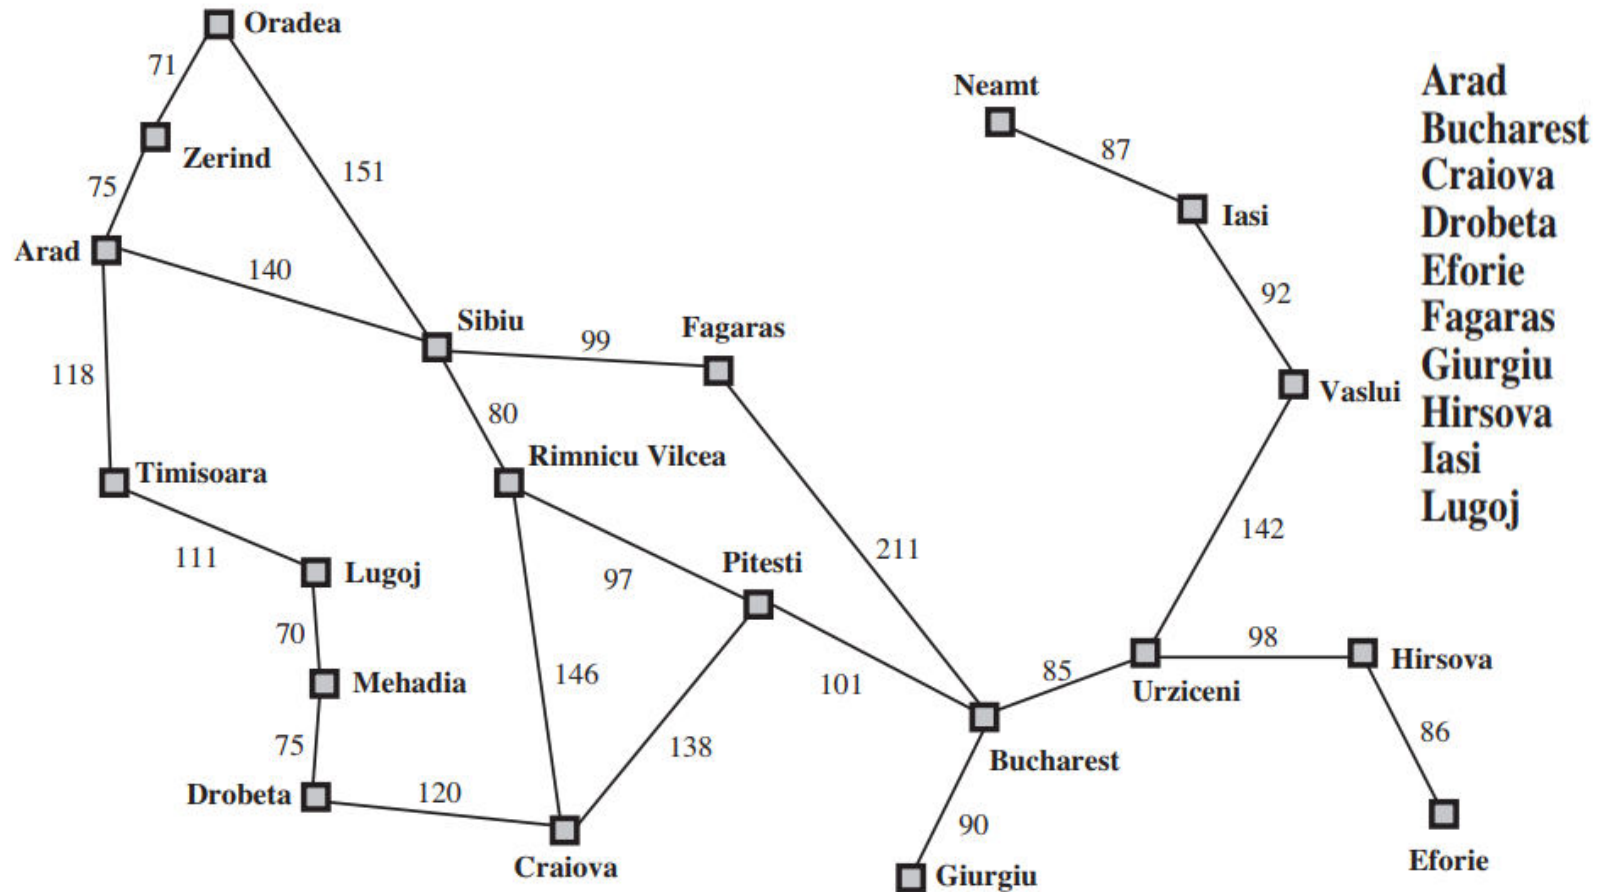
Informed (Heuristic) Search

- **Informed search** strategy—one that uses problem-specific knowledge beyond the definition of the problem itself
- Can find solutions more efficiently than can an uninformed strategy.
- A **heuristic function**, denoted $h(n)$:
 - $h(n)$ = estimated cost of the cheapest path from the state at node n to a goal state
 - Example: $h(\text{Arad})$ = the cost of the cheapest path from Arad to Bucharest via the straight-line distance from Arad to Bucharest (Goal).

Greedy best-first search

- Greedy best-first search tries to expand the node that is closest to the goal
- It evaluates nodes by using just the heuristic function; that is, $f(n) = h(n)$
 - Example: If the goal is Bucharest, $h_{SLD}(In(Arad)) = 366$, the straight line distance from Arad to the Bucharest.

Greedy best-first search Example




Arad	366
Bucharest	0
Craiova	160
Drobeta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244


Mehadia	241
Neamt	234
Oradea	380
Pitesti	100
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Greedy best-first search Example

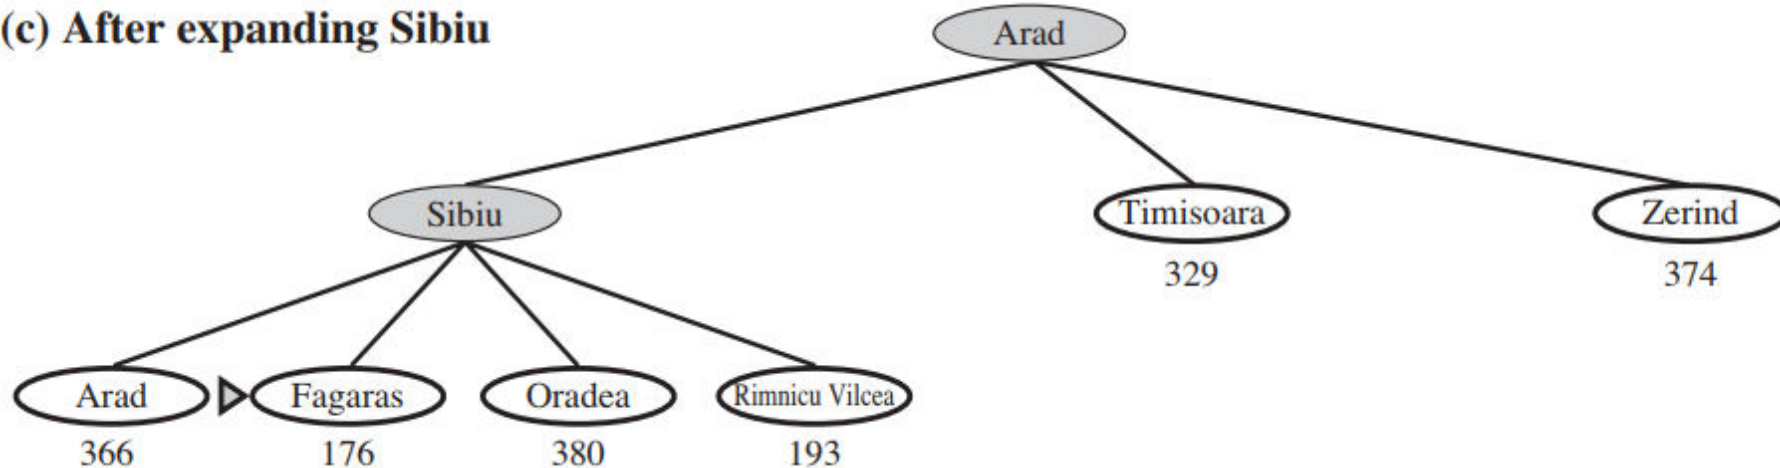
(a) The initial state



(b) After expanding Arad

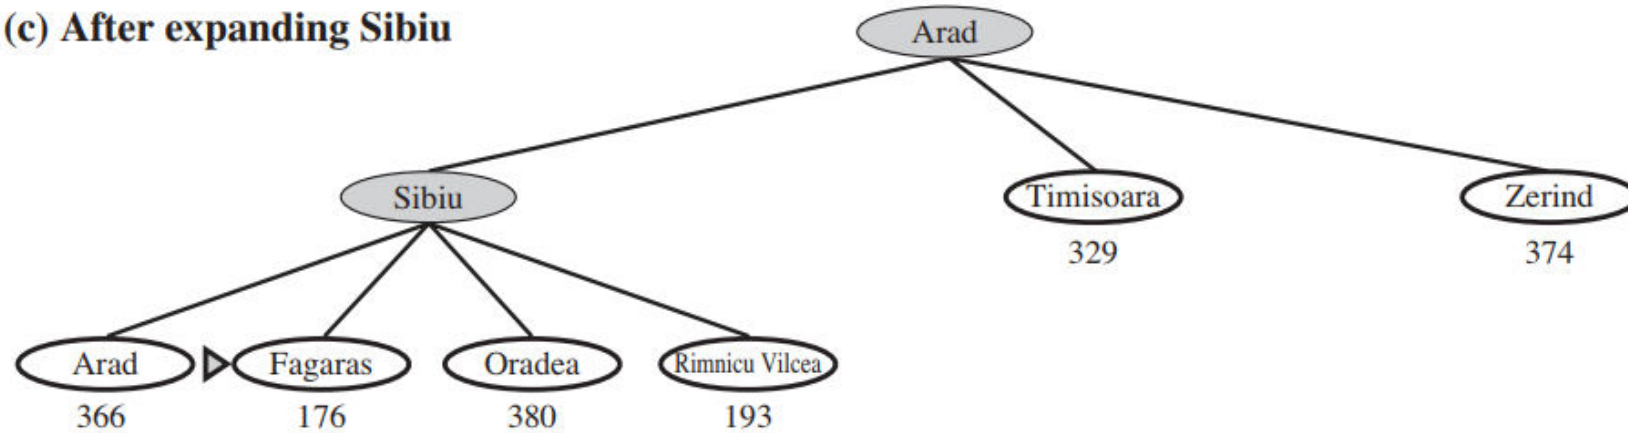


(c) After expanding Sibiu

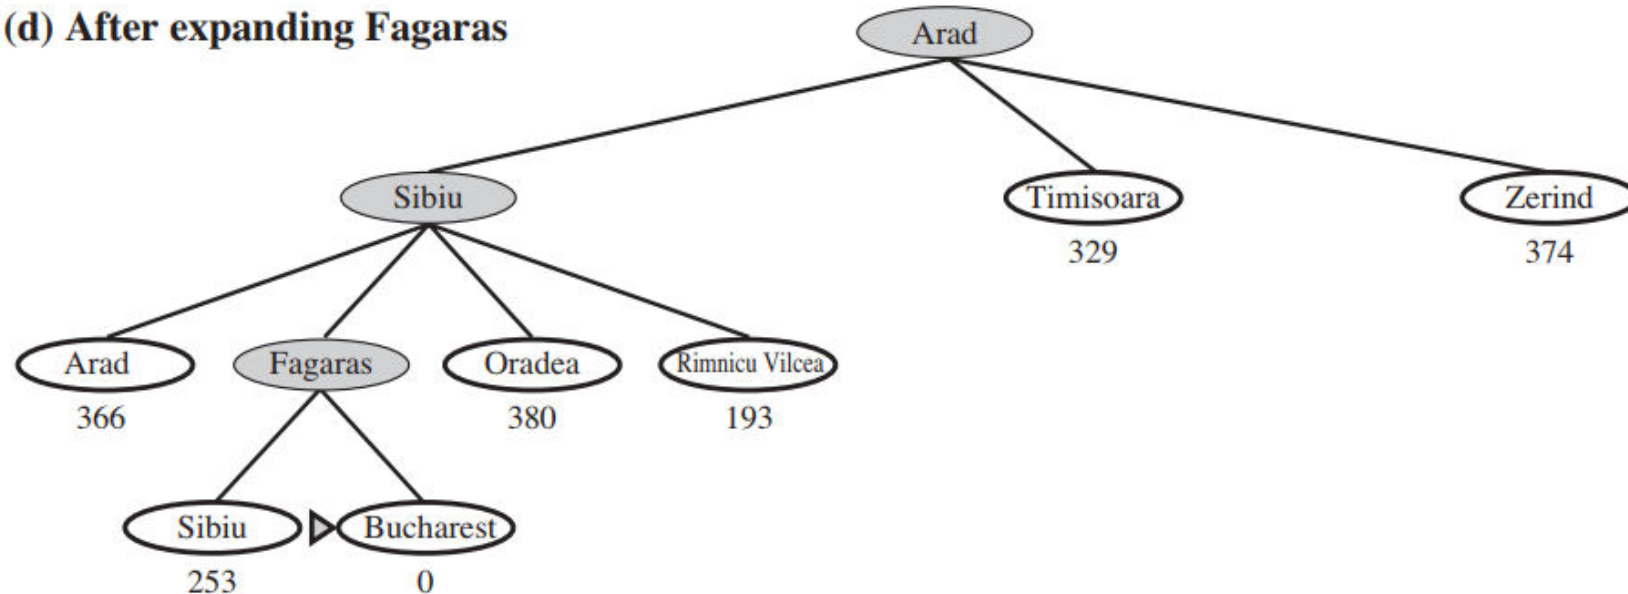


Greedy best-first search Example

(c) After expanding Sibiu



(d) After expanding Fagaras



A* Search

- The most widely known form of best-first search is called A* search
- the cost
- The estimated cost of the cheapest solution through n :

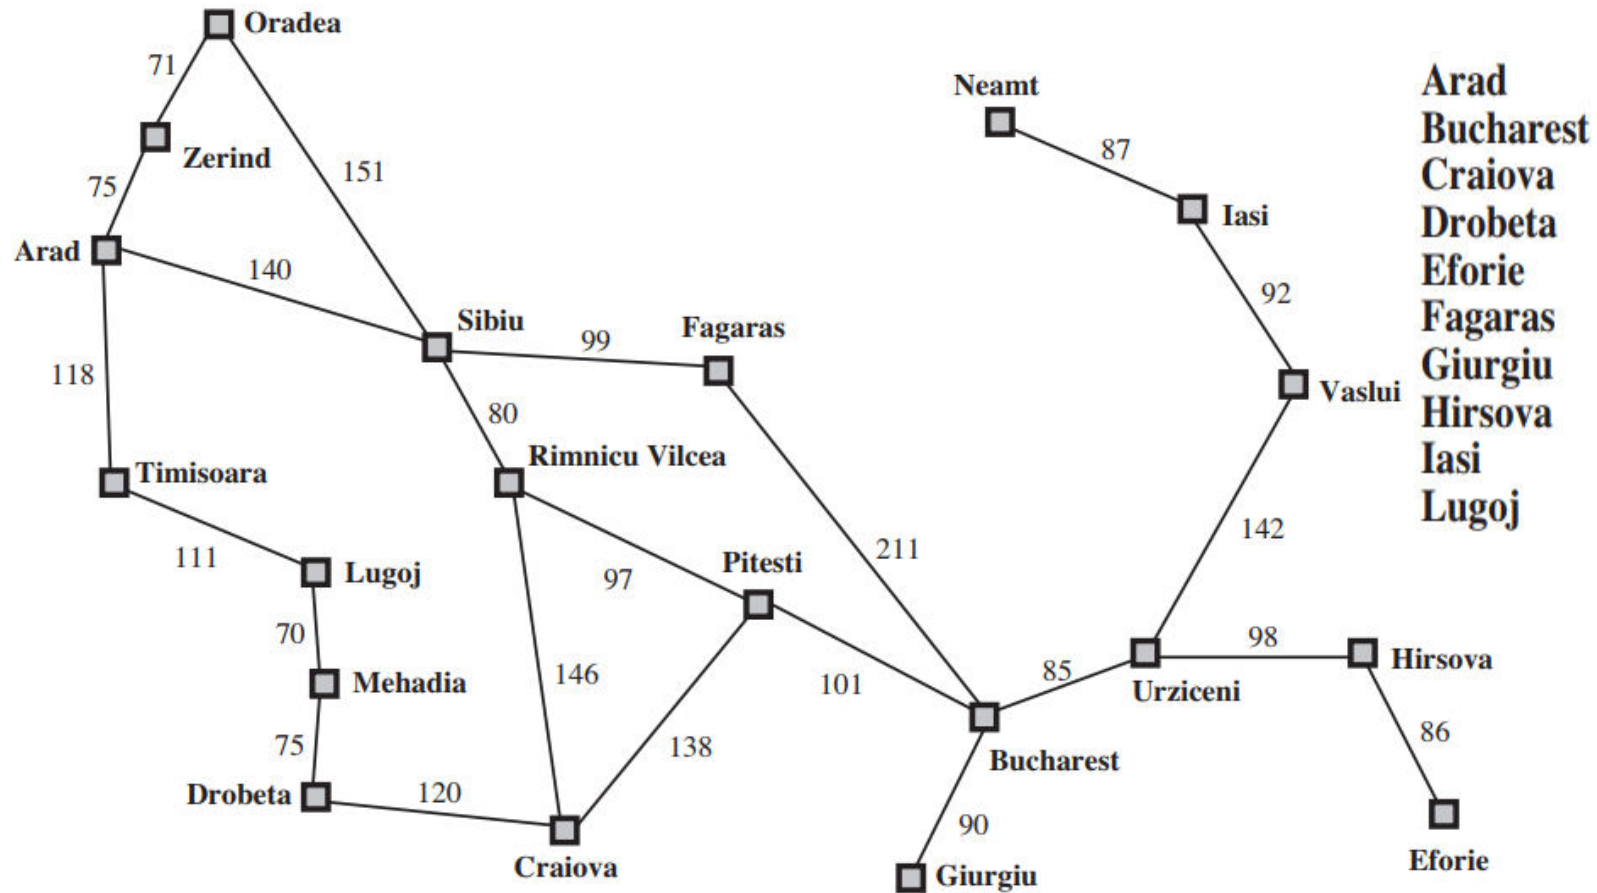
$$f(n) = g(n) + h(n)$$

- $g(n)$ gives the path cost from the *start node* to node n ,
 - $h(n)$ is the estimated cost of the cheapest path from n to the *goal*
- Algorithm uses two lists:
 - Frontier/open list ← the set of all leaf nodes available for expansion (can be maintained by a priority queue ordered by $f(n)$)
 - Explored/closed list ← the set of all explored nodes

A* Search Pseudocode

```
function A*-SEARCH(problem) returns a solution, or failure
node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE,  $f(\textit{node}) = g(\textit{node})$ 
frontier  $\leftarrow$  a priority queue ordered by  $f$ -cost, with node as the only element initially
explored  $\leftarrow$  an empty set
loop do
    if EMPTY?(frontier) then return failure
    node  $\leftarrow$  POP(frontier) /* chooses the lowest  $f$ -cost node in frontier */
    if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
        child  $\leftarrow$  CHILD-NODE(problem, node, action)
         $f(\textit{child}) = g(\textit{child}) + h(\textit{child})$ 
        if child.STATE is not in explored or frontier then
            frontier  $\leftarrow$  INSERT(child, frontier)
        else if child.STATE is in frontier with higher  $f(\textit{child})$  then
            replace that frontier node with child
```


A* Search Example




Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

A* Search Example


(a) The initial state



(b) After expanding Arad




(c) After expanding Sibiu

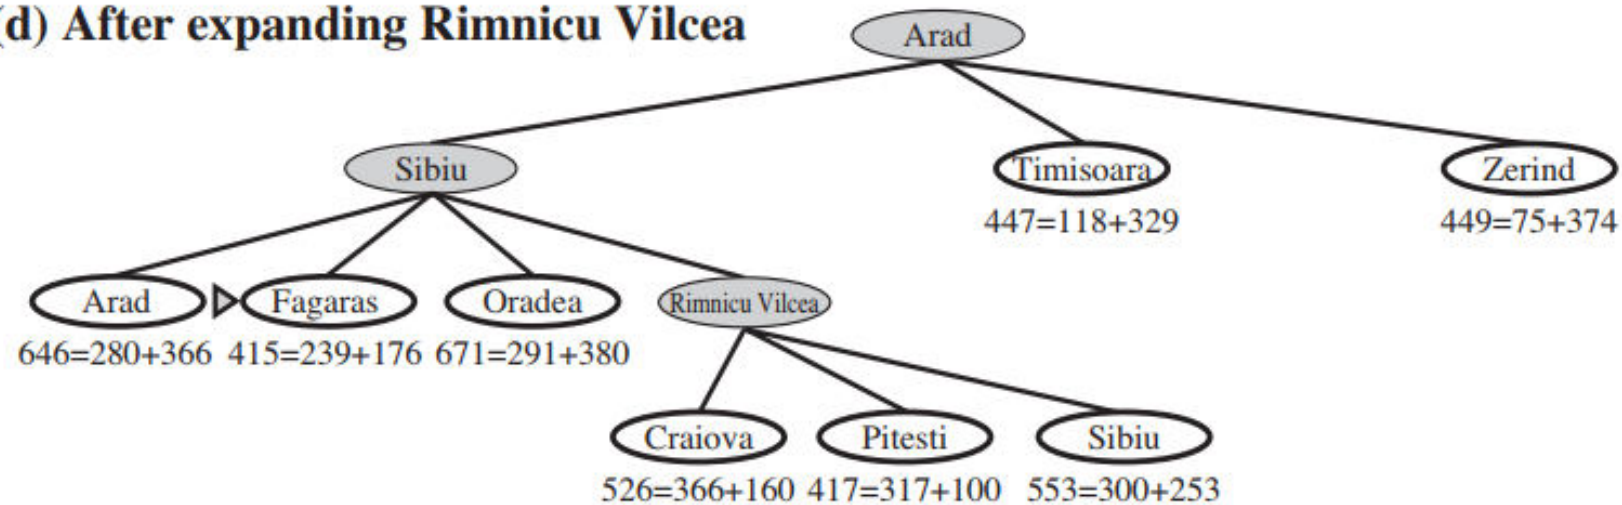


A* Search Example

(c) After expanding Sibiu

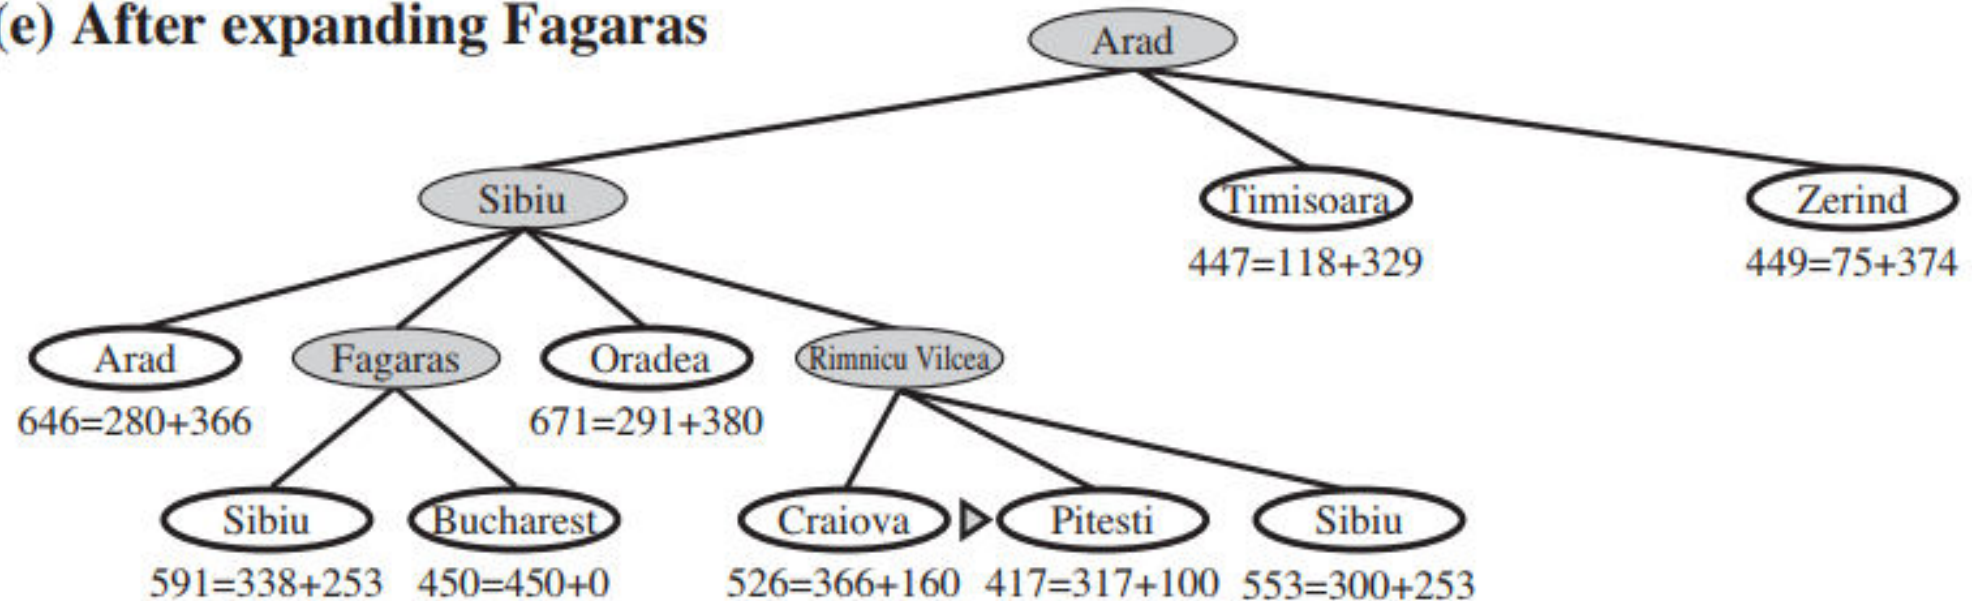


(d) After expanding Rimnicu Vilcea



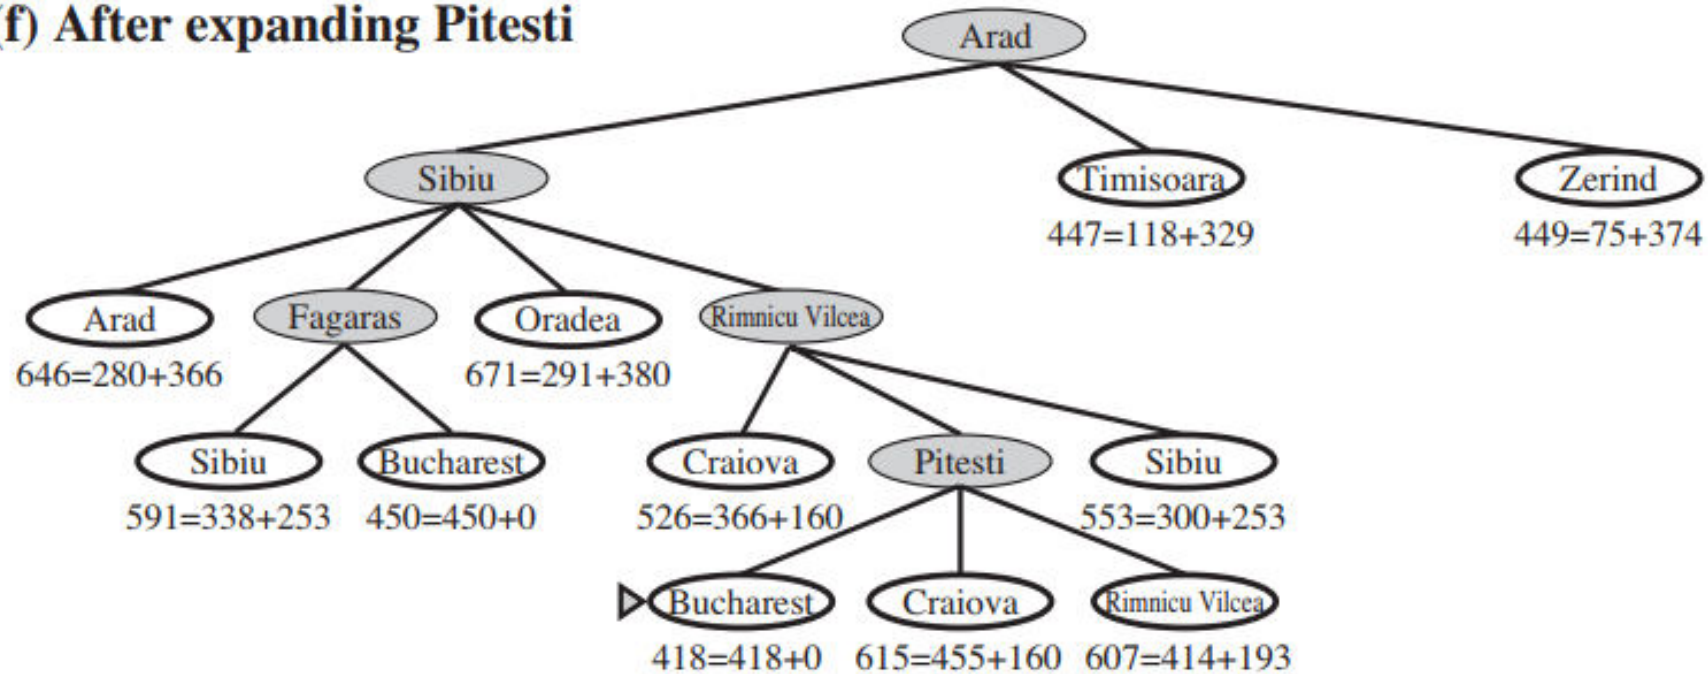
A* Search Example

(e) After expanding Fagaras



A* Search Example

(f) After expanding Pitesti




Admissible heuristics

- An admissible heuristic never overestimates the cost to reach the goal, i.e., it is optimistic.
- A heuristic $h(n)$ is admissible if for every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the true cost to reach the goal state from n .
- Example: $h_{SLD}(n) = \textit{Straight Line Distance}$ (never overestimates the actual road distance)

Optimality of A* Search

- If $h(n)$ is admissible, A* using TREE- SEARCH is optimal.
- Proof:
- Suppose some suboptimal goal $G2$ has been generated and is in the frontier. Let n be an unexpanded node in the frontier such that n is on a shortest path to an optimal goal G .




Optimality of A* Search

- $f(G2) = g(G2)$, since $h(G2) = 0$
- $f(G) = g(G)$, since $h(G) = 0$
- $g(G2) > g(G)$, since $G2$ is suboptimal
- Therefore, $f(G2) > f(G)$
- $h(n) \leq h^*(n)$, since h is admissible
- $g(n) + h(n) \leq g(n) + h^*(n)$
- $f(n) \leq g(n) + h^*(n) < f(G) < f(G2)$
- Hence $f(G2) > f(n)$, and A* will never select $G2$ for expansion.

Optimality of A* Search

- If $h(n)$ is consistent, A* using GRAPH-SEARCH is optimal.
- Proof:
- A heuristic is consistent (or monotonic) if for every node n , every successor n' of n generated by any action a : $h(n) \leq c(n, a, n') + h(n')$



Optimality of A* Search

- If h is consistent, we have:

$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n, a, n') + h(n') \\ &\geq g(n) + h(n) = f(n) \end{aligned}$$

- i.e., $f(n)$ is non-decreasing along any path.

Properties of A*

- Complete: Yes (unless there are infinitely many nodes with $f \leq f(G)$).
- Time: Exponential.
- Space: Keeps all nodes in memory, so also exponential.
- Optimal: Yes (provided h admissible or consistent).
- Optimally Efficient: Yes (no algorithm with the same heuristic is guaranteed to expand fewer nodes).